

Chapter 6: Synchronization

同步Tools





Chapter 6: Synchronization 同步 Tools

Background

The Critical-Section Problem

Peterson's Solution

Synchronization Hardware

Mutex Locks

Semaphores

Monitors





Objectives

To present the concept of process synchronization.

To introduce the critical-section problem, whose solutions can be used to ensure the consistency of shared data

To present both software and hardware solutions of the critical-section problem

To examine several classical process-synchronization problems

To explore several tools that are used to solve process synchronization problems





Background

Processes can execute concurrently

May be interrupted **at any time, partially completing** execution

Concurrent access to shared data may result in **data inconsistency**

Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes

Illustration of the problem:

Suppose that we wanted to provide a solution to the **consumer-producer problem** that fills **all** the buffers. We can do so by having an integer **counter** that keeps track of the number of full buffers. Initially, **counter** is set to 0. It is **incremented by the producer** after it produces a new buffer and is decremented by the consumer after it consumes a buffer.





Producer

```
while (true) {  
    /* produce an item in next produced */  
  
    while (counter == BUFFER_SIZE) ;  
        /* do nothing */  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
}
```





Consumer

```
while (true) {  
    while (counter == 0)  
        ; /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
    /* consume the item in next consumed */  
}
```





Race Condition

counter++ could be implemented as

```
register1 = counter
register1 = register1 + 1
counter = register1
```

counter-- could be implemented as

```
register2 = counter
register2 = register2 - 1
counter = register2
```

Consider this execution interleaving with “count = 5” initially:

S0: producer execute	<code>register1 = counter</code>	{register1 = 5}
S1: producer execute	<code>register1 = register1 + 1</code>	{register1 = 6}
S2: consumer execute	<code>register2 = counter</code>	{register2 = 5}
S3: consumer execute	<code>register2 = register2 - 1</code>	{register2 = 4}
S4: producer execute	<code>counter = register1</code>	{counter = 6}
S5: consumer execute	<code>counter = register2</code>	{counter = 4}

中間被分離掉

<https://ithelp.ithome.com.tw/articles/10225917>

goroutine 是輕量級執行緒(lightweight thread)





Critical Section Problem

Consider system of n processes $\{p_0, p_1, \dots, p_{n-1}\}$

Each process has **critical section** segment of code

Process **may be changing common variables**, updating table, writing file, etc

When **one process in critical section**, no other may be in its critical section

Critical section problem is to design **protocol** to solve this

Each process must **ask permission** to enter critical section in **entry section**, may follow critical section with **exit section**, then **remainder section**





Critical Section

General structure of process P_i

```
do {  
    entry section  
    critical section  
    exit section  
    remainder section  
} while (true);
```





Algorithm for Process P_i

```
do {  
    while (turn == j);  
        critical section  
    turn = j;  
        remainder section  
} while (true);
```





Solution to Critical-Section Problem

1. **Mutual Exclusion(排除)**- If process P_i is executing in its critical section, then **no other processes** can **be executing** in their critical sections
2. **Progress** - If no process is executing in its **critical section** and there exist some processes that wish to enter their critical section, then the selection of the processes that will enter the critical section next **cannot** be **postponed indefinitely(推遲無限期)**
3. **Bounded Waiting** - **A bound** must exist on the number of times that other processes are allowed to enter their critical sections after a process has **made a request** to enter its critical section and **before that request is granted(確認)**
 - Assume that each process executes at a **nonzero speed**
 - **No assumption concerning relative speed** of the n processes





Critical-Section Handling in OS

Two approaches depending on if **kernel** is preemptive or non-preemptive

Preemptive – allows **preemption of process** when running in **kernel mode**

Non-preemptive – **runs until exits kernel** mode, blocks, or **voluntarily**(自願) yields CPU

▶ **Essentially** (基本) **free** of race conditions in kernel mode





Peterson's Solution

Good algorithmic description of solving the problem

Two process solution

Assume that the **load** and **store** machine-language instructions are **atomic**; that is, **cannot be interrupted**

The two processes share two variables:

```
int turn;
```

```
Boolean flag[2]
```

The variable `turn` indicates **whose turn** it is to **enter** the critical section

The **flag** array is used to **indicate** if **a process is ready to enter** the critical section. **flag[i] = true** implies that process **P_i** is ready!





Algorithm for Process P_i

do {

```
flag[i] = true;
```

```
turn = j;
```

```
while (flag[j] && turn == j);
```

critical section

```
flag[i] = false;
```

remainder section

```
} while (true);
```





Peterson's Solution (Cont.)

Provable 可證明 that the three CS requirement are met:

1. **Mutual exclusion is preserved**

P_i enters CS only if:

either `flag[j] = false` or `turn = i`

2. **Progress requirement is satisfied**
3. **Bounded-waiting requirement is met**





Synchronization Hardware

Many systems provide hardware support for implementing the critical section code.

All solutions below based on idea of **locking**

Protecting **critical regions** via locks

Uniprocessors – could disable interrupts

Currently running code would execute without preemption

Generally too inefficient (低效率) on multiprocessor systems

- ▶ Operating systems using this not broadly 寬廣 scalable

Modern machines provide special **atomic** hardware instructions(指示)

- ▶ **Atomic** = non-interruptible(原子同一時間只能有一個執行緒對共享資源進行讀寫操作)

Either test memory word and set value

Or **swap contents** of two memory words





Solution to Critical-section Problem Using Locks

```
do {  
    acquire lock  
        critical section  
    release lock  
    remainder section  
} while (TRUE);
```





test_and_set Instruction

Definition:

```
boolean test_and_set (boolean *target)
{
    boolean rv = *target;
    *target = TRUE;
    return rv;
}
```

1. Executed atomically
2. Returns the original value of passed parameter
3. Set the new value of passed parameter to “TRUE”.





Solution using test_and_set()

Shared Boolean variable lock, initialized to FALSE

Solution:

```
do {  
    while (test_and_set(&lock))  
        ; /* do nothing */  
        /* critical section */  
    lock = false;  
        /* remainder section */  
} while (true);
```





compare_and_swap Instruction

Definition:

```
int compare_and_swap(int *value, int expected, int new_value) {  
    int temp = *value;  
  
    if (*value == expected)  
        *value = new_value;  
  
    return temp;  
}
```

1. Executed atomically
2. **Returns** the original value of passed parameter “value”
3. Set the variable “value” the value of the passed parameter “**new_value**” but only if “value” == “**expected**”. That is, the swap takes place only under this condition.





Solution using compare_and_swap

Shared integer “lock” initialized to 0;

Solution:

```
do {  
    while (compare_and_swap(&lock, 0, 1) != 0)  
        ; /* do nothing */  
    /* critical section */  
    lock = 0;  
    /* remainder section */  
} while (true);
```





Bounded-waiting Mutual Exclusion with test_and_set

```
do {
    waiting[i] = true;
    key = true;
    while (waiting[i] && key)
        key = test_and_set(&lock);
    waiting[i] = false;
    /* critical section */
    j = (i + 1) % n;
    while ((j != i) && !waiting[j])
        j = (j + 1) % n;
    if (j == i)
        lock = false;
    else
        waiting[j] = false;
    /* remainder section */
} while (true);
```





Mutex Locks

Previous solutions are complicated and generally **inaccessible** to application programmers

OS designers build software tools to solve critical section problem

Simplest is mutex lock

Protect a critical section by first **acquire()** a lock then **release()** the lock

Boolean variable indicating if lock is available or not

Calls to **acquire()** and **release()** must be atomic

Usually **implemented via hardware atomic instructions**

But this solution requires **busy waiting**

This lock therefore called a **spinlock** 自旋鎖





acquire() and release()

```
acquire() {  
    while (!available)  
        ; /* busy wait */  
    available = false;;  
}  
  
release() {  
    available = true;  
}  
  
do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
} while (true);
```





Semaphore信號

Synchronization tool that provides more **sophisticated** (複雜的) ways (than Mutex locks) for process to synchronize their activities.

Semaphore **S** – integer variable

Can only be accessed via **two indivisible** 不可分割 (atomic) operations

wait() and **signal()**

- ▶ Originally called **P()** and **V()**

Definition of the **wait()** operation

```
wait(S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}
```

Definition of the **signal()** operation

```
signal(S) {  
    S++;  
}
```





Semaphore Usage

Counting semaphore – integer value can range over an unrestricted無限制的 domain

Binary semaphore – integer value can range only between 0 and 1

Same as a **mutex lock (互斥)**

Can solve various synchronization problems

Consider P_1 and P_2 that require S_1 to happen before S_2

Create a semaphore “synch” initialized to 0

P1 :

S_1 ;

signal(synch) ;

P2 :

wait(synch) ;

S_2 ;

Can implement a counting semaphore S as a binary semaphore





Semaphore Implementation

Must guarantee that no two processes can execute the `wait()` and `signal()` on the same semaphore at the same time

Thus, the implementation becomes the critical section problem where the `wait` and `signal` code are placed in the critical section

Could now have **busy waiting** in critical section implementation

- ▶ But **implementation code is short**
- ▶ **Little busy** waiting if critical section rarely occupied (critical section 狀況少發生)

Note that **applications** may spend lots of time in critical sections and therefore this is not a good solution





Semaphore Implementation with no Busy waiting

With each semaphore there is an associated waiting queue

Each entry in a waiting queue has two data items:

value (of type integer)

pointer to next record in the list

Two operations:

block – place the process invoking(請求) the operation on the **appropriate waiting queue**

wakeup – remove one of processes in the **waiting queue** and **place it in the ready queue**

```
typedef struct{  
    int value;  
    struct process *list;  
} semaphore;
```





Implementation with no Busy waiting (Cont.)

```
wait(semaphore *S) {
    S->value--;
    if (S->value < 0) {
        add this process to S->list;
        block();
    }
}

signal(semaphore *S) {
    S->value++;
    if (S->value <= 0) {
        remove a process P from S->list;
        wakeup(P);
    }
}
```





Deadlock and Starvation

Deadlock – two or more processes are **waiting indefinitely** for an event that can be **caused** by **only one** of the waiting processes

Let S and Q be two semaphores initialized to 1

P_0
`wait(S);`
`wait(Q);`
`...`
`signal(S);`
`signal(Q);`

P_1
`wait(Q);`
`wait(S);`
`...`
`signal(Q);`
`signal(S);`

Starvation 飢餓 – **indefinite(無限)blocking**

A process may never be removed from the semaphore queue in which it is suspended

Priority Inversion – Scheduling problem when lower-priority process holds a lock needed by **higher-priority** process

Solved via **priority-inheritance(繼承) protocol**





Problems with Semaphores

Incorrect use of semaphore operations:

signal (mutex) wait (mutex)

wait (mutex) ... wait (mutex)

Omitting (省略) of wait (mutex) or signal (mutex) (or both)

Deadlock and starvation are possible.





Monitors

A high-level abstraction that provides a convenient and effective mechanism for process synchronization

Abstract data type, internal variables (内部) only accessible by code within the procedure

Only one process may be active within the monitor at a time

But not powerful enough to model some synchronization schemes

```
monitor monitor-name
{
    // shared variable declarations
    procedure P1 (...) { ... }

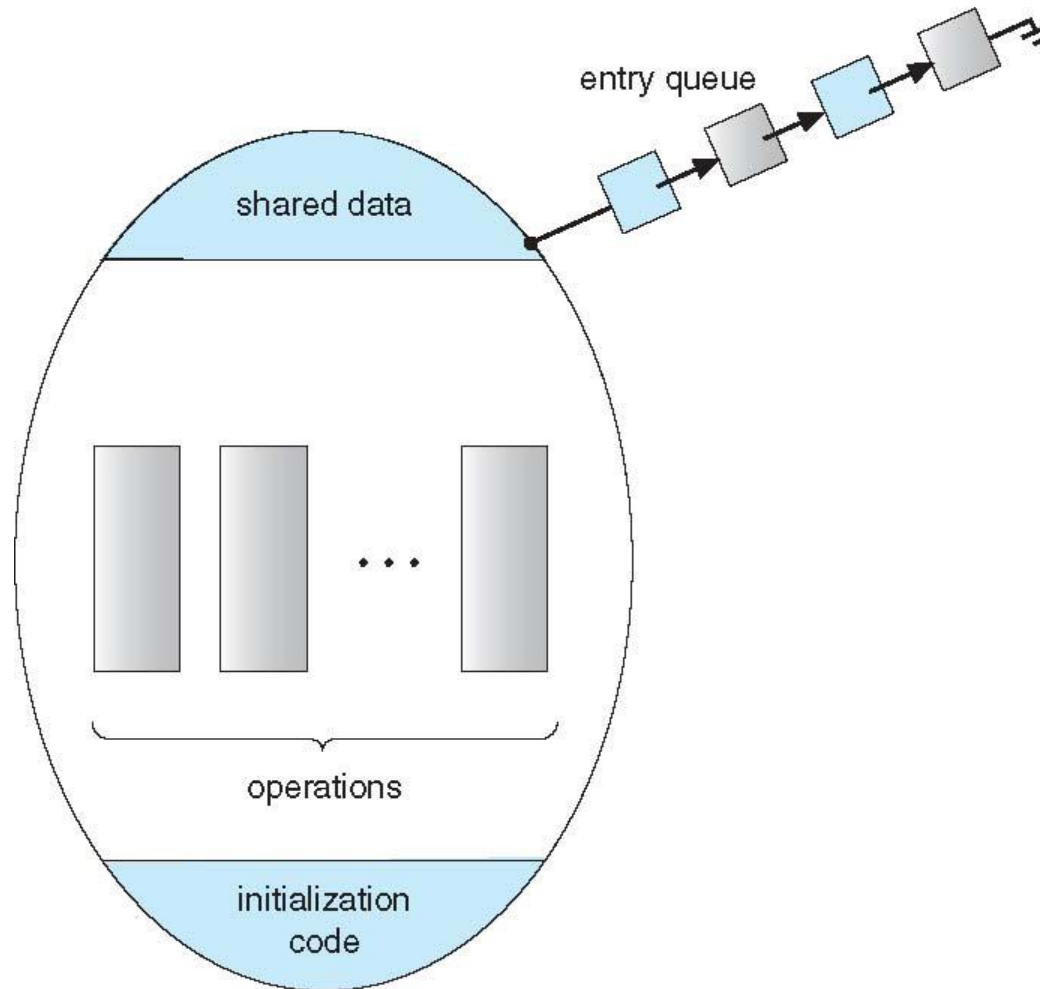
    procedure Pn (...) {.....}

    Initialization code (...) { ... }
}
}
```





Schematic view of a Monitor





Condition Variables

condition **x**, **y**;

Two operations are allowed on a condition variable:

x.wait() – a process that invokes the operation is suspended until **x.signal()**

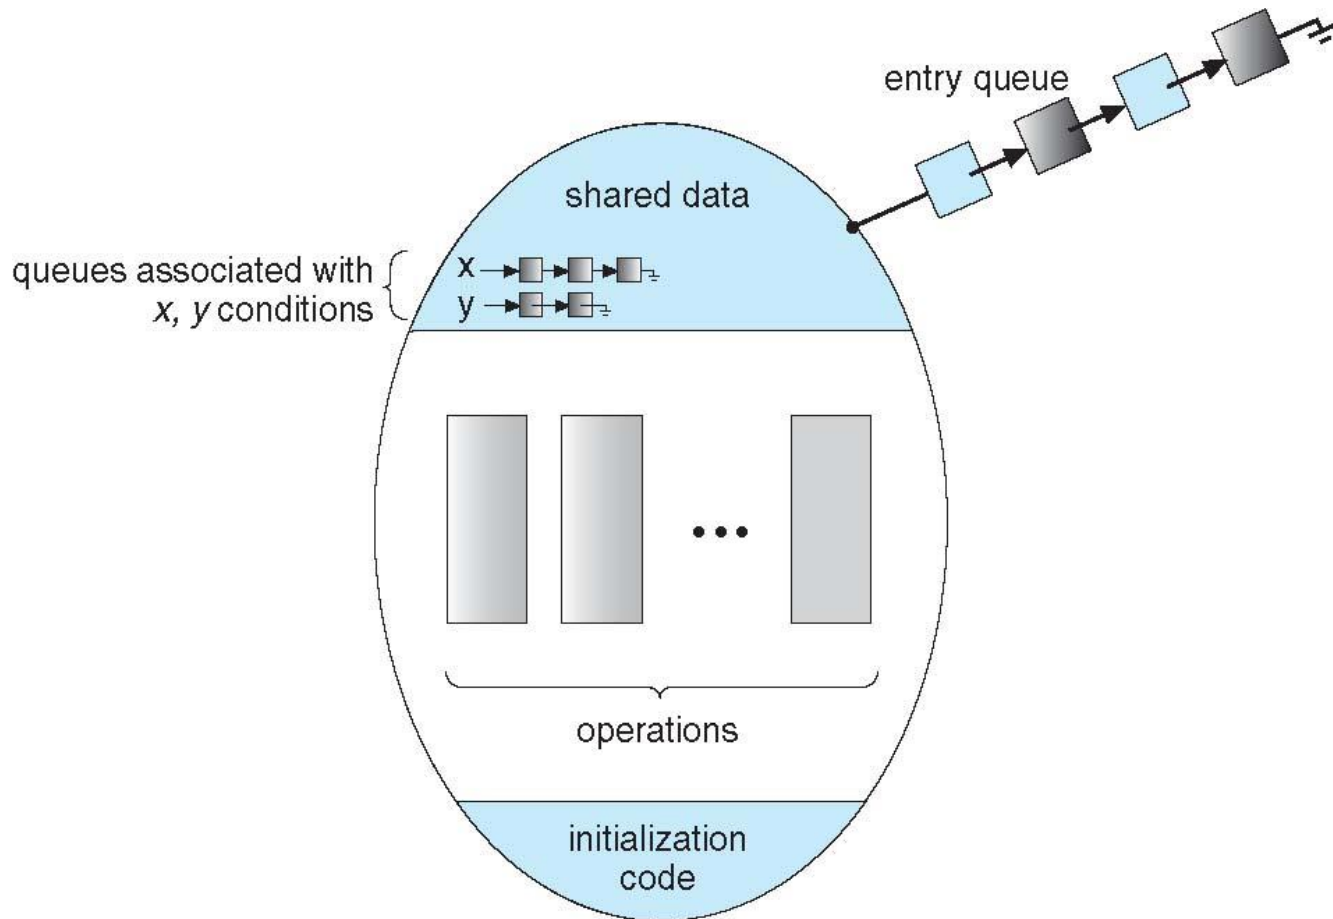
x.signal() – resumes one of processes (if any) that invoked **x.wait()**

- ▶ If no **x.wait()** on the variable, then it has no effect on the variable





Monitor with Condition Variables





Condition Variables Choices

If process P invokes `x.signal()`, and process Q is suspended in `x.wait()`, what should happen next?

Both Q and P cannot execute in parallel. If Q is resumed, then P must wait

Options include

Signal and wait – P waits until Q either leaves the monitor or it waits for another condition

Signal and continue – Q waits until P either leaves the monitor or it waits for another condition

Both have pros and cons – language implementer can decide

Monitors implemented in Concurrent **Pascal** compromise

- ▶ P executing signal immediately leaves the monitor, Q is resumed

Implemented in other languages including Mesa, **C#, Java**





Monitor Implementation Using Semaphores

Variables

```
semaphore mutex;  // (initially = 1)
semaphore next;   // (initially = 0)
int next_count = 0;
```

Each procedure F will be replaced by

```
wait(mutex) ;
...
body of F;
...
if (next_count > 0)
    signal(next)
else
    signal(mutex) ;
```

Mutual exclusion within a monitor is ensured





Monitor Implementation – Condition Variables

For each condition variable x , we have:

```
semaphore x_sem; // (initially = 0)
int x_count = 0;
```

The operation $x.\text{wait}$ can be implemented as:

```
x_count++;
if (next_count > 0)
    signal(next);
else
    signal(mutex);
wait(x_sem);
x_count--;
```





Monitor Implementation (Cont.)

The operation `x.signal` can be implemented as:

```
if (x_count > 0) {  
    next_count++;  
    signal(x_sem);  
    wait(next);  
    next_count--;  
}
```





Resuming恢復Processes within a Monitor

If several processes queued on condition x, and x.signal() **executed**, which should be **resumed**?

FCFS frequently not adequate充足

conditional-wait construct of the form **x.wait(c)**

Where c is **priority number**

Process with **lowest number (highest priority)** is scheduled next





Single Resource allocation

Allocate a single resource among competing processes using priority numbers that specify the maximum time a process plans to use the resource

```
R.acquire (t) ;  
    ...  
    access the resource ;  
    ...  
  
R.release ;
```

Where R is an instance of type ResourceAllocator





A Monitor to Allocate **Single** Resource

```
monitor ResourceAllocator
{
    boolean busy;
    condition x;
    void acquire(int time) {
        if (busy)
            x.wait(time);
        busy = TRUE;
    }
    void release() {
        busy = FALSE;
        x.signal();
    }
    initialization code() {
        busy = FALSE;
    }
}
```



End of Chapter 6

