

# Chapter 7: Synchronization Examples

---





# Chapter 7: Synchronization Examples

---

## Classic Problems of Synchronization

Synchronization within the Kernel

POSIX Synchronization

Synchronization in Java

Alternative Approaches





# Classical Problems of Synchronization

---

Classical problems used to test newly-proposed synchronization schemes

Bounded-Buffer Problem

Readers and Writers Problem

Dining-Philosophers Problem





# Bounded-Buffer Problem

---

***n*** buffers, each can hold **one item**

Semaphore (訊號) **mutex** initialized to the **value 1**

Semaphore **full** initialized to the value 0

Semaphore **empty** initialized to the value **n**





# Bounded Buffer Problem (Cont.)

The structure of the **producer process**

```
do {  
    ...  
    /* produce an item in next_produced */  
    ...  
    wait(empty);  
    wait(mutex);  
    ...  
    /* add next produced to the buffer */  
    ...  
    signal(mutex);  
    signal(full);  
} while (true);
```





# Bounded Buffer Problem (Cont.)

The structure of the **consumer** process

```
Do {  
    wait(full) ;  
    wait(mutex) ;  
    ...  
    /* remove an item from buffer to next_consumed */  
    ...  
    signal(mutex) ;  
    signal(empty) ;  
    ...  
    /* consume the item in next consumed */  
    ...  
} while (true) ;
```





# Readers-Writers Problem

A data set is shared among a number of concurrent processes

Readers – only read the data set; they do **not** perform any updates

Writers – can both read and write

Problem – allow multiple readers to read at the same time

Only one single writer can access the shared data at the same time

Several variations of how readers and writers are considered – all involve some form of priorities優先順序

## Shared Data

Data set

Semaphore **rw\_mutex** initialized to 1 (讀寫)

Semaphore **mutex** initialized to 1 (讀)

Integer **read\_count** initialized to 0





# Readers-Writers Problem (Cont.)

---

The structure of a writer process

```
do {  
    wait(rw_mutex);  
    ...  
    /* writing is performed */  
    ...  
    signal(rw_mutex);  
} while (true);
```







# Readers-Writers Problem (Cont.)

The structure of a reader process

```
do {  
    wait(mutex);  
    read_count++;  
    if (read_count == 1)  
        wait(rw_mutex);  
    signal(mutex);  
  
    ...  
    /* reading is performed */  
    ...  
    wait(mutex);  
    read_count--;  
    if (read_count == 0)  
        signal(rw_mutex);  
    signal(mutex);  
} while (true);
```





# Readers-Writers Problem Variations

**First** variation變化 – no reader kept waiting unless **writer** has permission to use shared object (除非有writer去使用共享object, 否則不能有reader在一直等待)

**Second** variation – once writer is ready, it performs the write ASAP (as soon as possible)


Both may have starvation飢餓 leading to even more variations(上述都可能產生飢餓情況)

Problem is solved on some systems by kernel providing (reader-writer locks: 解決上述的問題)





# Dining-Philosophers Problem



Philosophers **spend their lives alternating (一生中)** thinking and eating  
Don't interact with their **neighbors**, occasionally(偶爾) try to **pick up 2**  
chopsticks (one at a time) to **eat from bowl**

Need **both** to eat, then **release both** when done

In the case of 5 philosophers

**Shared data**

- ▶ Bowl of **rice** (data set)
- ▶ **Semaphore chopstick [5]** initialized to 1





# Dining-Philosophers Problem Algorithm

---

Deadlock的發生：一開始大家都拿右方的筷子並沒有問題，但當大家要拿取左方的筷子時，會因為已經被坐在自己左邊的人給取走了，所以就會形成一個「Deadlock」的problem狀態。

Deadlock的處理：

只允許在一個圓桌上有4個人而已。

要檢查左右是否都已經available，可以的話便能eating。

使用asymmetric方法解決：也就是說，坐在奇位數位置的人先拿取左邊再拿取右邊的筷子；坐在雙數位置的人先拿取右邊再拿取左邊的筷子，方能解決問題。

。





# Dining-Philosophers Problem Algorithm

The structure of Philosopher *i*:

```
do {  
    wait (chopstick[i] );  
    wait (chopStick[ (i + 1) % 5] );  
  
    // eat  
    signal (chopstick[i] );  
    signal (chopstick[ (i + 1) % 5] );  
  
    // think  
} while (TRUE);
```

What is the **problem with this algorithm?**

**signal(mutex).....wait(mutex) / wait(mutex).....wait(mutex)  
會遺漏wait(mutex)或是signal(mutex), 亦或是兩個一起遺漏掉。**





# Monitor Solution to Dining Philosophers

```
monitor DiningPhilosophers
{
    enum { THINKING; HUNGRY, EATING} state [5] ;
    condition self [5];

    void pickup (int i) {
        state[i] = HUNGRY;
        test(i);
        if (state[i] != EATING) self[i].wait;
    }

    void putdown (int i) {
        state[i] = THINKING;
        // test left and right neighbors
        test((i + 4) % 5);
        test((i + 1) % 5);
    }
}
```





# Solution to Dining Philosophers (Cont.)

```
void test (int i) {
    if ((state[(i + 4) % 5] != EATING) &&
        (state[i] == HUNGRY) &&
        (state[(i + 1) % 5] != EATING) ) {
        state[i] = EATING ;
        self[i].signal () ;
    }
}

initialization_code() {
    for (int i = 0; i < 5; i++)
        state[i] = THINKING;
}
}
```





# Solution to Dining Philosophers (Cont.)

Each philosopher  $i$  invokes the operations **pickup()** and **putdown()** in the following sequence:

```
DiningPhilosophers.pickup(i);
```

**EAT**

```
DiningPhilosophers.putdown(i);
```

每個philosopher  $i$  invoke兩個operations **pickup()** 以及 **putdown()** 在sequence中。

因為一次一個進入, 所以並不會產生dealock, 但因為動作慢所以可能會產生starvation的情況。

No deadlock, but starvation is possible







# A Monitor to Allocate Single Resource

```
monitor ResourceAllocator
{
    boolean busy;
    condition x;
    void acquire(int time) {
        if (busy)
            x.wait(time);
        busy = TRUE;
    }
    void release() {
        busy = FALSE;
        x.signal();
    }
    initialization code() {
        busy = FALSE;
    }
}
```

Monitor:提供方便且有效率的機制給process synchronization使用。有些系統不支持。一次只能有一個process可以在monitor內執行。不能解決所有的synchronization問題，但這已經是較為複雜的方法





# Synchronization Examples

---

Solaris

Windows

Linux

Pthreads





# Solaris Synchronization

Implements a variety of locks to support **multitasking**, **multithreading** (including **real-time threads**), and **multiprocessing**

Uses **adaptive mutexes** for efficiency when protecting data from short code segments

Starts as a standard semaphore **spin-lock**(標準的**semaphore spin-lock**)

If **lock** held, and by a **thread running** on another CPU, spins(保持**lock**，以及被**thread**執行在其他的CPU,spins)

If lock held by **non-run-state** thread, **block and sleep** waiting for **signal** of lock being **released** (**lock**被不在執行狀態的**thread**給**block**和**sleep**，等待單一的**lock**被釋放)

Uses **condition variables**

Uses **readers-writers locks** when longer sections of code need access to data (當code需要longer section去存取data)





# Solaris Synchronization(2)

Uses **turnstiles** to order the list of threads waiting to acquire either an adaptive mutex or reader-writer lock (使用turnstiles去命令在互相等待的threads，將其整理成一個list，等待mutex或是reader-writer解決，並以資料結構的方式處理。)

Turnstiles are **per-lock-holding-thread**, not per-object

**Priority**-inheritance per-turnstile gives the **running thread** the **highest** of the priorities of the threads in its turnstile (將priority-inheritance一個turnstile給正在執行thread最高優先全的threads在這個turnstile旋轉門)





# Windows Synchronization

Uses **interrupt masks** to protect access to **global resources** on **uniprocessor** systems(保護global resources在uniprocessor systems存取)

Uses **spinlocks** on **multiprocessor systems**

Spinlocking-thread will **never be preempted**(永不會被中斷)

Also provides **dispatcher objects** user-land (使用物件分配器給使用者空間) which may **act** **mutexes**, **semaphores**, events, and timers

## Events

- ▶ An **event** acts much like a **condition variable**(類似於 condition variable)

**Timers** notify one or more thread when **time expired** (timer 會通知一個或多個thread)

Dispatcher objects either **signaled-state** (object **available**) or **non-signaled state** (thread will **block**)





# Linux Synchronization

---

Linux:

Prior to **kernel Version 2.6**, **disables** interrupts to implement **short critical sections** (但回應速度較慢)

Version 2.6 and later, fully preemptive (回應速度較佳)

Linux provides:

**Semaphores**

**atomic integers**

**spinlocks**

**reader-writer** versions of both

On single-cpu system, **spinlocks** **replaced** by enabling and disabling kernel preemption (在單一CPU系統時, spinlock可被可用或不可用的kernel preemption所取代)





# Pthreads Synchronization

---

Pthreads API is OS-independent

It provides:

mutex locks

condition variable

Non-portable (可移植) extensions(擴增) include:

read-write locks

spinlocks





# Alternative(其他) Approaches

---

Transactional Memory

OpenMP

Functional Programming Languages







# Transactional Memory

A **memory transaction (交易)** is a **sequence of read-write operations to memory that are performed atomically (不被中斷)**.

```
void update()  
{  
    /* read/write memory */  
}
```





# OpenMP

OpenMP is a set of **compiler** directives and API(compiler指令與API所組成) that support **parallel programming**支持平行程式執行.

```
void update(int value)
{
    #pragma omp critical
    {
        count += value
    }
}
```

The code contained within the **#pragma omp critical** directive is treated as a critical section and performed atomically(在程式碼中有#pragma omp critical 的指令, 就像對待CS一樣是個atomic, 不會中斷。).





# Functional Programming Languages

Functional programming languages offer a different paradigm(範例) than procedural languages in that they do not maintain state(在不能維持的狀態).

Variables are treated as immutable(不變) and cannot change state once they have been assigned a value.

There is increasing interest in functional languages such as Erlang and Scala for their approach in handling data races(處理 data race condition).



# End of Chapter 7

---

